





# Lesson 6:

## Sentence Classification and a little about RNNs

Partially based on slides by Jurafsky and Martin  
Speech and Language Processing, 3<sup>rd</sup> Edition

# Sentiment Analysis

## *Example #1: Movie Reviews*

-  • Unbelievably disappointing
-  • Full of zany characters and richly applied satire, and some great plot twists
-  • This is the greatest screwball comedy ever filmed
-  • It was pathetic. The worst part about it was the boxing scenes.

# Sentiment Analysis

## Example #2: Product Reviews



**HP Officejet 6500A Plus e-All-in-One Color Ink-jet - Fax / copier / printer / scanner**

**\$89 online, \$100 nearby** ★★★★★ 377 reviews

September 2010 - Printer - HP - Inkjet - Office - Copier - Color - Scanner - Fax - 250 sheets

### Reviews

**Summary** - Based on 377 reviews




What people are saying

ease of use	<div><div></div></div>	"This was very easy to setup to four computers."
value	<div><div></div></div>	"Appreciate good quality at a fair price."
setup	<div><div></div></div>	"Overall pretty easy setup."
customer service	<div><div></div></div>	"I DO like honest tech support people."
size	<div><div></div></div>	"Pretty Paper weight."
mode	<div><div></div></div>	"Photos were fair on the high quality mode."
colors	<div><div></div></div>	"Full color prints came out with great quality."

- A Sentiment Analysis system called CALM applied to Twitter predicts the Dow Jones Industrial Average (DJIA) 3 days later
- Such algorithms are already used by hedge funds

Johan Bollen, Huina Mao, Xiaojun Zeng. 2011. [Twitter mood predicts the stock market](#), Journal of Computational Science 2:1, 1-8. 10.1016/j.jocs.2010.12.007.



# Scherer Typology of Affective States

- **Emotion:** brief organically synchronized ... evaluation of a major event
  - *angry, sad, joyful, fearful, ashamed, proud, elated*
- **Mood:** diffuse non-caused low-intensity long-duration change in subjective feeling
  - *cheerful, gloomy, irritable, listless, depressed, buoyant*
- **Interpersonal stances:** affective stance toward another person in a specific interaction
  - *friendly, flirtatious, distant, cold, warm, supportive, contemptuous*
- **Attitudes:** enduring, affectively colored beliefs, dispositions towards objects or persons
  - *liking, loving, hating, valuing, desiring*
- **Personality traits:** stable personality dispositions and typical behavior tendencies
  - *nervous, anxious, reckless, morose, hostile, jealous*

# Sentiment Analysis: Definition

- Sentiment analysis is the detection of **attitudes**  
“enduring, affectively colored beliefs, dispositions towards objects or persons”
  1. **Holder (source)** of attitude
  2. **Target (aspect)** of attitude
  3. **Type** of attitude
    - From a set of types: *like, love, hate, value, desire*, etc.
    - Or (more commonly) simple weighted **polarity**: *positive, negative, neutral*, together with *strength*
  4. **Text** containing the attitude
    - Sentence or entire document

# Sentiment Analysis

- Simplest task:
  - Is the attitude of this text positive or negative?
- More complex:
  - Rank the attitude of this text from 1 to 5
- Advanced:
  - Detect the target, source, or complex attitude types

# Sentiment Classification in Movie Reviews

- Is an IMDB movie review positive or negative?



when \_star wars\_ came out some twenty years ago  
, the image of traveling throughout the stars has  
become a commonplace image . [...]

when han solo goes light speed , the stars change  
to bright lines , going towards the viewer in lines  
that converge at an invisible point .

cool .



“ snake eyes ” is the most aggravating  
kind of movie : the kind that shows so  
much potential then becomes  
unbelievably disappointing .

it's not just because this is a brian  
depalma film , and since he's a great  
director and one who's films are always  
greeted with at least some fanfare .

Bo Pang, Lillian Lee, and Shivakumar Vaithyanathan. 2002. Thumbs up? Sentiment Classification using Machine Learning Techniques. EMNLP-2002, 79—86.

Bo Pang and Lillian Lee. 2004. A Sentimental Education: Sentiment Analysis Using Subjectivity Summarization Based on Minimum Cuts. ACL, 271-278



# A Simple Classifier

- Log-linear or Naïve Bayes classifier
- Features:
  - Tokenized words
  - Possibly mark-up (e.g., hashtags in Twitter, headers in HTMLs)
- Features are often binary
  - Indicating whether a word appeared or did not appear in the document (bag of words)
  - Often works better for text classification than word frequency

# Error Analysis:

## What makes reviews hard to classify?

- **Sarcasm:**

- Perfume review in *Perfumes: the Guide*:
  - “If you are reading this because it is your darling fragrance, please wear it at home exclusively, and tape the windows shut.”
- On *Automobile Steering Wheel Attachable Work Surface*:
  - “You wouldn’t believe how much more interesting my commute is now that I have something to do other than just stare out the window! I’m using it right now to post this review and I never”



- **Thwarted Expectations and Ordering Effects:**

- “This film should be brilliant. It sounds like a great plot, the actors are first grade, and the supporting cast is good as well, and Stallone is attempting to deliver a good performance. However, it can’t hold up.”

# Negation in Sentiment Analysis

- One practice is to add NOT\_ to every word between negation and following punctuation:

didn't like this movie , but I



didn't NOT\_like NOT\_this NOT\_movie but I

# Negation in Sentiment Analysis

- This is a very crude solution:
  - Explicit negation is only one way to reverse polarity
    - “He **failed** to convey the importance of his message”
  - Logical structure can be more complex
    - “I don’t think anyone could have done a better job”
  - Negation scope:
    - “I didn’t like the exposition, but otherwise liked the movie”
- More recent approaches take the context (surrounding words) into account

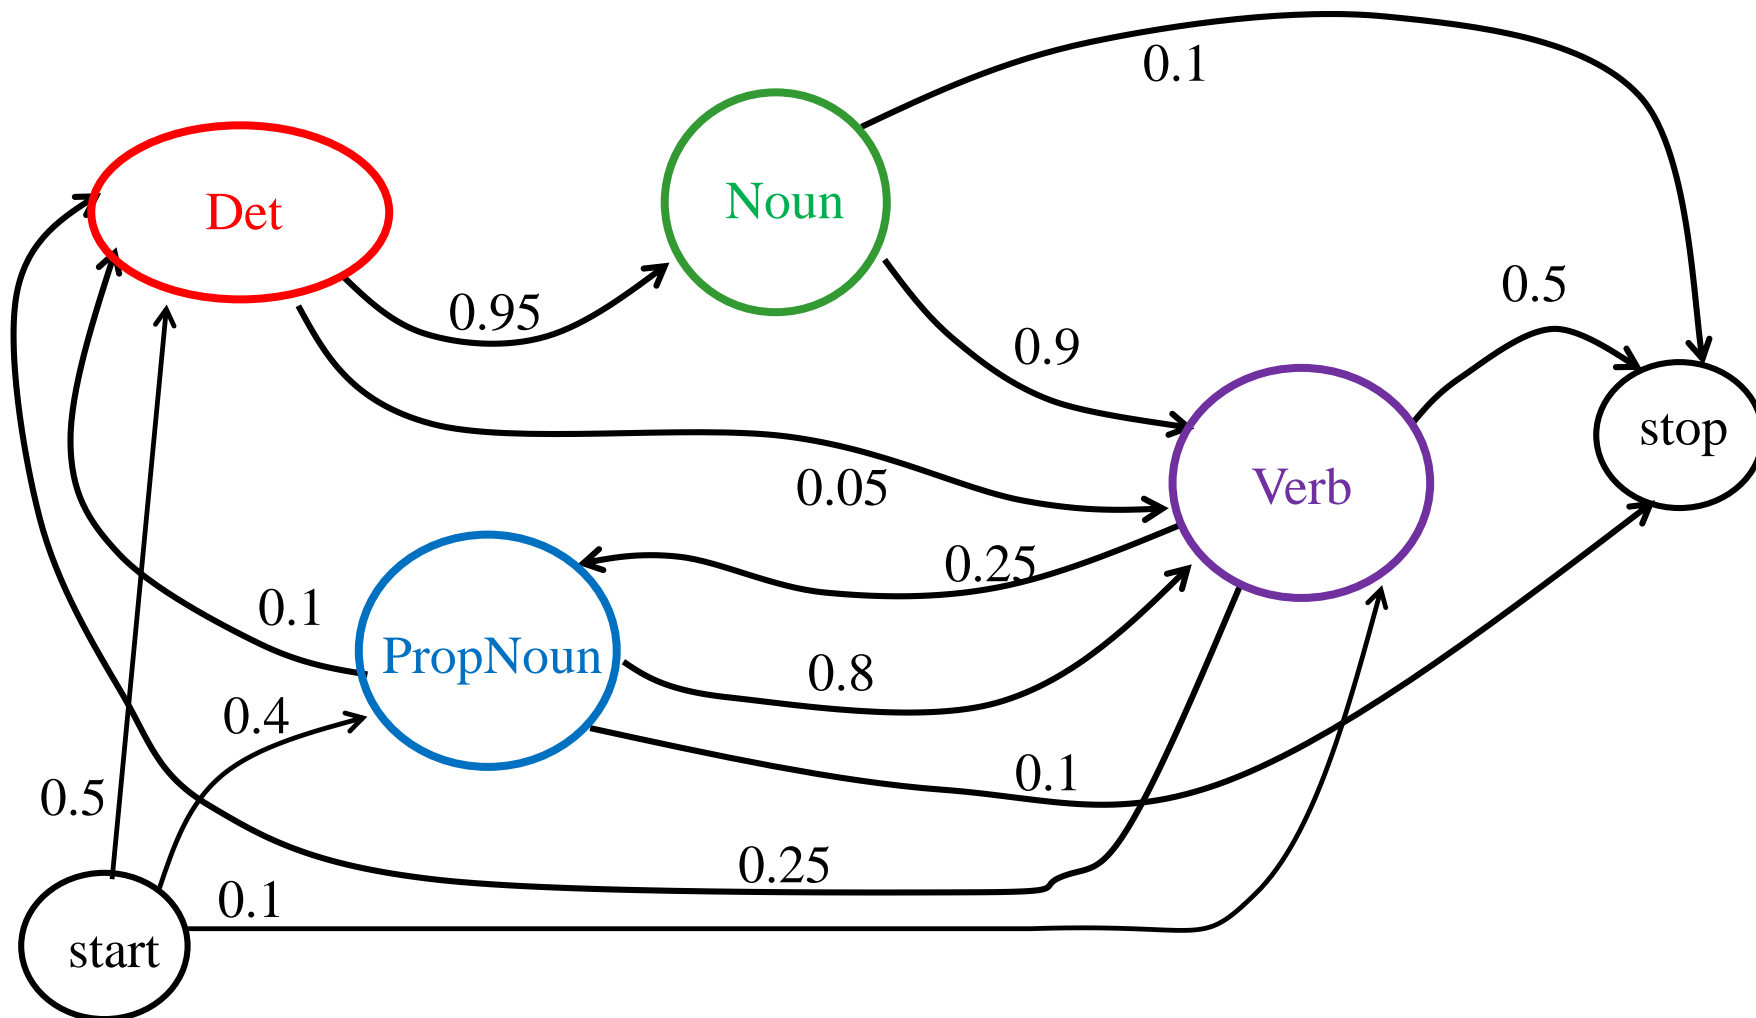
# Sentiment Analysis as Sequence Labeling

- This construal of sentiment analysis attempts to capture the meaning of a word in context by encoding (parts of) the sentence as features
- Recently: using Recurrent Neural Networks (RNNs)
- Recall the underlying assumption in Markov (n-gram) models:
  - You only need to know the last  $n$  tokens you've encountered to know what's next
  - Alternative view: the probability of a sequence is the product of the probabilities of its *n-grams*

# Sentiment Analysis as Sequence Labeling

- Consider the example:
  - *How can you not see this movie?*
  - *You should not see this movie*
- How well will a bi-gram model work?
  - Similar unigrams and bigrams → similar prediction
- The problem with Markov Models: need to maintain a **state** to capture distant influences
  - The size of the space increases exponentially with the distance

Recall: Markov Models are FSA with transition probabilities



# Recurrent Neural Networks

- Motivation:
  - Neural network model, but with a state
  - How can we borrow ideas from FSAs?
- RNNs are FSAs ...
  - With a twist
  - No longer finite in the same sense
  - The state is an embedding of the history in a continuous space



# Recurrent Neural Networks

- Map from dense sequence to dense representation
  - Maps a sequence of vector inputs to a sequence of vector states

$$x_1, \dots, x_n \rightarrow s_1, \dots, s_n$$

- A (parametrized) transition function  $R$  does the mapping:

$$s_i = R(s_{i-1}, x_i)$$

- $R$  is parametrized and parameters are shared across steps

$$s_4 = R(s_3, x_4) = R(R(s_2, x_3), x_4) = R(R(R(s_1, x_2), x_3), x_4)$$

# Recurrent Neural Networks

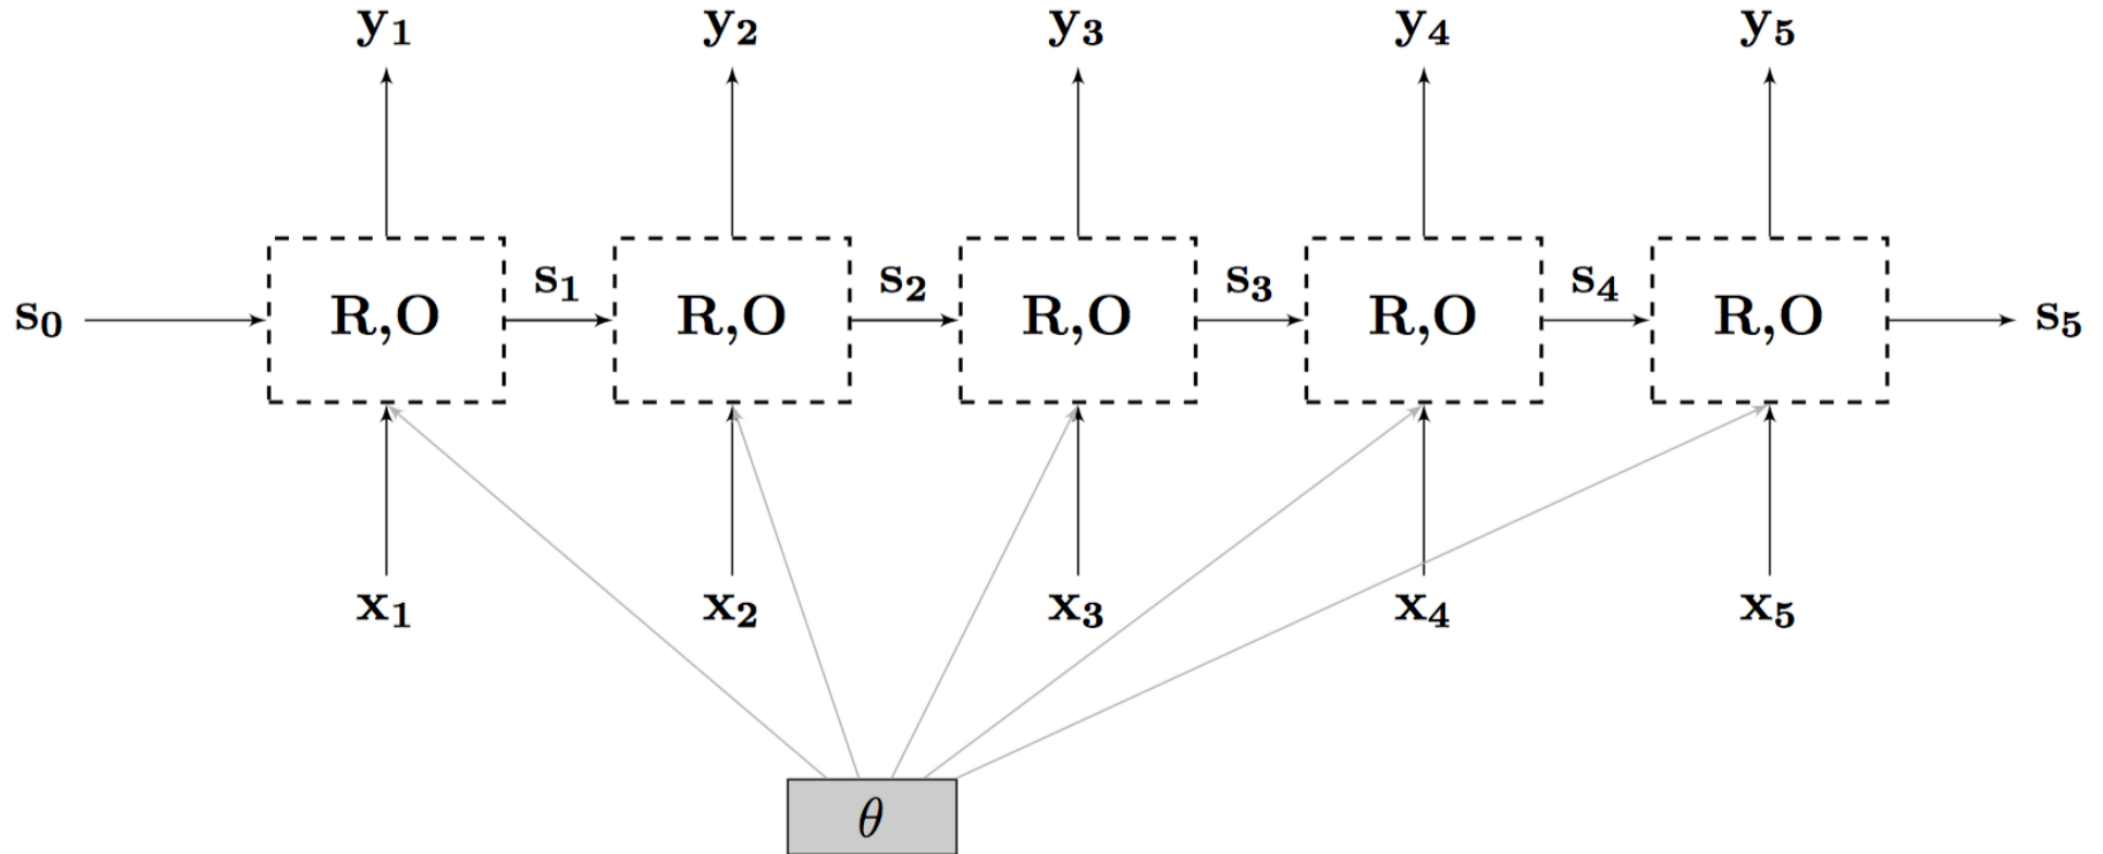
- An output function  $O$  maps states to (vector) outputs, which are often viewed as a distribution over the possible labels

- Example:

$$R_{Elman}(s, x) = \tanh(W[x, s] + b)$$

$$O(s_i) = \text{softmax}(W s_i + b)$$

# RNNs: Graphical Representation




# Back to Sentiment Analysis

- When using RNNs for sentence classification (such as sentiment analysis), it is often practical to use Bi-RNNs
- Bi-RNNs:
  - 2 RNNs, one going back to forth and the other forth to back
  - Output function is a function of both states
- This allows the states associated with each word to encode relevant information from the words following them and preceding them

$$\vec{h}_t = \mathcal{H}(W_{x\vec{h}}x_t + W_{\vec{h}\vec{h}}\vec{h}_{t-1} + b_{\vec{h}}) \quad (9)$$

$$\overleftarrow{h}_t = \mathcal{H}(W_{x\overleftarrow{h}}x_t + W_{\overleftarrow{h}\overleftarrow{h}}\overleftarrow{h}_{t+1} + b_{\overleftarrow{h}}) \quad (10)$$

$$y_t = W_{\vec{h}y}\vec{h}_t + W_{\overleftarrow{h}y}\overleftarrow{h}_t + b_y \quad (11)$$



# Back to Sentiment Analysis

- One simple way to do sentiment analysis (or other sentence classification) with Bi-RNNs is to average the output sequence:

$$y = \frac{1}{N} \sum_i y_i$$


- Now train a binary logistic classifier for predicting the sentiment:

$$P(+|y) = \frac{1}{1 + e^{w^t y + b}}$$

$$\vec{h}_t = \mathcal{H}(W_{x\vec{h}} x_t + W_{\vec{h}\vec{h}} \vec{h}_{t-1} + b_{\vec{h}}) \quad (9)$$

$$\overleftarrow{h}_t = \mathcal{H}(W_{x\overleftarrow{h}} x_t + W_{\overleftarrow{h}\overleftarrow{h}} \overleftarrow{h}_{t+1} + b_{\overleftarrow{h}}) \quad (10)$$

$$y_t = W_{\vec{h}y} \vec{h}_t + W_{\overleftarrow{h}y} \overleftarrow{h}_t + b_y \quad (11)$$

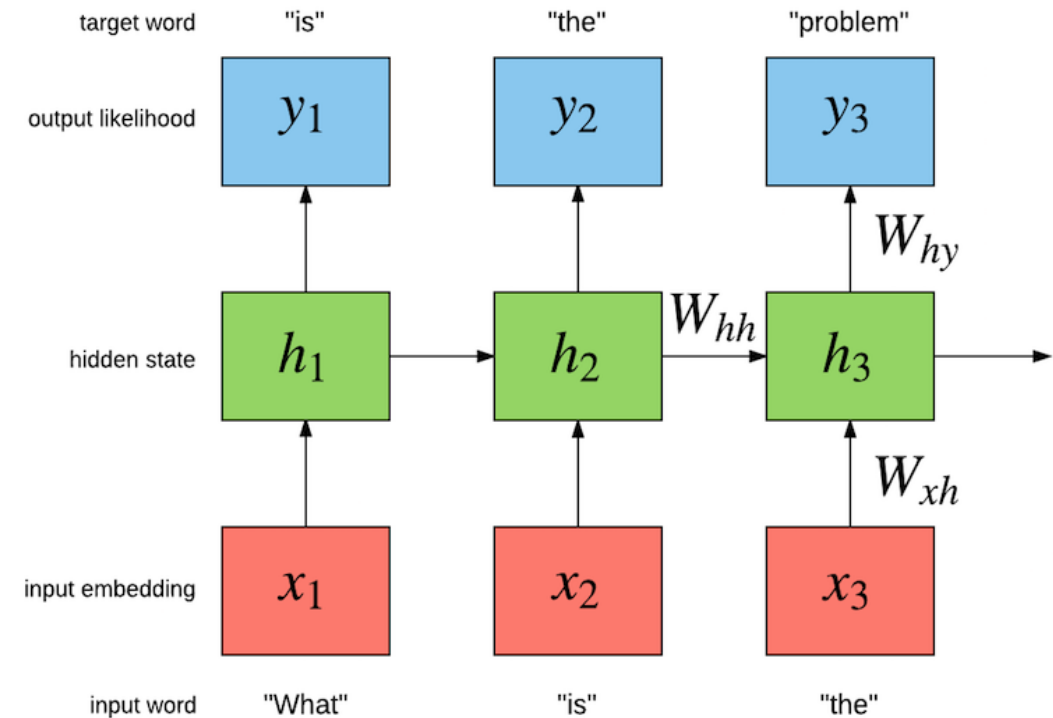


# Context in RNN models for Sentence Classification

- RNN models are able to take the context (both preceding and following) into account, as well as the linear order between the words
  - bag-of-words models cannot
- They indeed show better performance in tasks such as sentiment analysis
- However, the state sequence is not easy to interpret
- Further investigation is needed to establish what contextual and semantic aspects of sentences are captured using these techniques

# Neural Language Models

- Language models based on RNNs have shown much power in recent years, consistently surpassing n-gram models
- The basic architecture is that of a sequence recurrent neural net (RNN)
  - Input words are converted to 1-hot vectors
  - The output is passed through a softmax layer, which defines the probability of the next word
  - The loss function is often just the log probability of the next word predicted by the model



# Character-level Language Models

- Vocabulary: characters instead of words
- Advantage:
  - Small vocabulary → compact model
  - Can generalize over morphologically similar words
- However:
  - Need to learn how to spell
  - Longer range dependencies between tokens



# Character-level Language Models

