

Transition-based parsing

Daniel Hershcovich

Natural Language Processing, 67658
December 24, 2017

Overview

- 1 Introduction
- 2 Transition systems
- 3 Greedy transition-based parsing
- 4 Dealing with error propagation

Introduction

Dependency parsing

Given sentence $w = (w_1, \dots, w_n)$, let $V_w = \{0, 1, \dots, n\}$ (the root node has index 0).

Derive dependency tree $T = (V_w, A)$ by finding the set of arcs $A \subset V_w \times \mathcal{L} \times V_w$, where \mathcal{L} is the set of possible edge labels. Equivalently—for each i , find w_i 's head and dependency label.

Dependency parsing

Given sentence $w = (w_1, \dots, w_n)$, let $V_w = \{0, 1, \dots, n\}$ (the root node has index 0).

Derive dependency tree $T = (V_w, A)$ by finding the set of arcs $A \subset V_w \times \mathcal{L} \times V_w$, where \mathcal{L} is the set of possible edge labels. Equivalently—for each i , find w_i 's head and dependency label.




$$V_w = \{0, 1, 2, 3\}, \quad A = \{(0, \text{root}, 2), (2, \text{nsubj}, 1), (2, \text{dobj}, 3)\}$$

Dependency parsing

Given sentence $w = (w_1, \dots, w_n)$, let $V_w = \{0, 1, \dots, n\}$ (the root node has index 0).

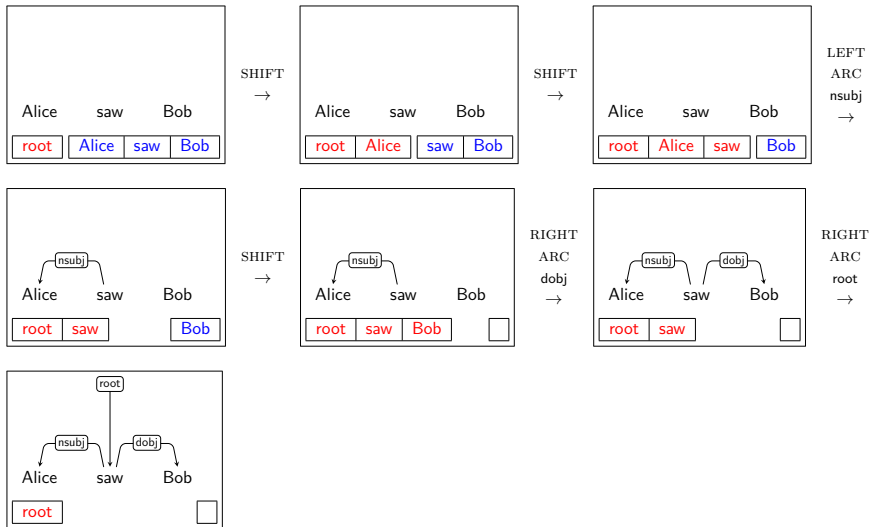
Derive dependency tree $T = (V_w, A)$ by finding the set of arcs $A \subset V_w \times \mathcal{L} \times V_w$, where \mathcal{L} is the set of possible edge labels. Equivalently—for each i , find w_i 's head and dependency label.



$$V_w = \{0, 1, 2, 3\}, \quad A = \{(0, \text{root}, 2), (2, \text{nsubj}, 1), (2, \text{dobj}, 3)\}$$

In transition-based parsing, the problem is decomposed to finding a sequence of *transitions*.

Example



Transition systems

Configurations

Transitions operate on the parser *configuration* (or *state*)

$$c = (\Sigma, B, A)$$

where

- $\Sigma \subseteq V_w$ is the stack of partially processed items.
- $B \subseteq V_w$ is the buffer of remaining input tokens.
- $A \subset V_w \times \mathcal{L} \times V_w$ is the set of arcs constructed so far.

Configurations

Transitions operate on the parser *configuration* (or *state*)

$$c = (\Sigma, B, A)$$

where

- $\Sigma \subseteq V_w$ is the stack of partially processed items.
- $B \subseteq V_w$ is the buffer of remaining input tokens.
- $A \subset V_w \times \mathcal{L} \times V_w$ is the set of arcs constructed so far.

Common notation:

$$\Sigma = [\dots, s_1, s_0] = \Sigma' | s_1 | s_0$$

$$B = [b_0, b_1, \dots] = b_0 | b_1 | B'$$

Transition systems

A *transition system* is defined as

$$S = (\mathcal{C}, \mathcal{T}, c_s, C_t)$$

where

- \mathcal{C} is the set of possible configurations.
- $\mathcal{T} \subset \mathcal{C}^2$ is the set of *transitions*.
- c_s maps every sentence w to an initial configuration $c_s(w)$.
- $C_t \subset \mathcal{C}$ is the set of terminal configurations.

Transition sequence

A *transition sequence* is $(c_0, \dots, c_m) \subseteq \mathcal{C}$ s.t.

- $c_0 = c_s(w)$
- $c_m = (\Sigma_m, B_m, A_m) \in C_t$
- For each $i = 1, \dots, m$ there exists $t \in \mathcal{T}$ s.t. $c_{i+1} = t(c_i)$.

The output of the system is then $T = (V_w, A_m)$.

Arc-standard transition system (Nivre, 2004)

Transition set \mathcal{T} :

SHIFT move one item from the buffer to the stack:

$$(\Sigma, i|B, A) \Rightarrow (\Sigma|i, B, A)$$

LEFT-ARC $_{\ell}$ create arc $s_0 \rightarrow s_1$ with label $\ell \in \mathcal{L}$ and remove s_1 :

$$(\Sigma|i|j, B, A) \Rightarrow (\Sigma|j, B, A \cup \{(j, \ell, i)\})$$

Condition: $i \neq 0$

RIGHT-ARC $_{\ell}$ create arc $s_1 \rightarrow s_0$ with label $\ell \in \mathcal{L}$ and remove s_0 :

$$(\Sigma|i|j, B, A) \Rightarrow (\Sigma|i, B, A \cup \{(i, \ell, j)\})$$

Arc-standard transition system (Nivre, 2004)

Transition set \mathcal{T} :

SHIFT move one item from the buffer to the stack:

$$(\Sigma, i|B, A) \Rightarrow (\Sigma|i, B, A)$$

LEFT-ARC $_{\ell}$ create arc $s_0 \rightarrow s_1$ with label $\ell \in \mathcal{L}$ and remove s_1 :

$$(\Sigma|i|j, B, A) \Rightarrow (\Sigma|j, B, A \cup \{(j, \ell, i)\})$$

Condition: $i \neq 0$

RIGHT-ARC $_{\ell}$ create arc $s_1 \rightarrow s_0$ with label $\ell \in \mathcal{L}$ and remove s_0 :

$$(\Sigma|i|j, B, A) \Rightarrow (\Sigma|i, B, A \cup \{(i, \ell, j)\})$$

Typically $|\mathcal{L}| \approx 50$, so there are 101 different transitions.

Arc-standard transition system (Nivre, 2004)

Transition set \mathcal{T} :

SHIFT move one item from the buffer to the stack:

$$(\Sigma, i|B, A) \Rightarrow (\Sigma|i, B, A)$$

LEFT-ARC $_{\ell}$ create arc $s_0 \rightarrow s_1$ with label $\ell \in \mathcal{L}$ and remove s_1 :

$$(\Sigma|i|j, B, A) \Rightarrow (\Sigma|j, B, A \cup \{(j, \ell, i)\})$$

Condition: $i \neq 0$

RIGHT-ARC $_{\ell}$ create arc $s_1 \rightarrow s_0$ with label $\ell \in \mathcal{L}$ and remove s_0 :

$$(\Sigma|i|j, B, A) \Rightarrow (\Sigma|i, B, A \cup \{(i, \ell, j)\})$$

Typically $|\mathcal{L}| \approx 50$, so there are 101 different transitions.

Initial configuration:

$$c_s(w_1, w_2, w_3, \dots) = ([0], [1, 2, 3, \dots], \emptyset)$$

Terminal configuration:

$$c_t = ([0], [], A)$$

Properties of the arc-standard system

Soundness. Every transition sequence outputs a projective tree.

Completeness. Every projective tree is output by some sequence.


Complexity. Input of length n requires exactly $2n$ transitions.

Bottom-up. Attaches a token's head only after all dependents.


¹<http://aclweb.org/anthology/J08-4003>

Projectivity

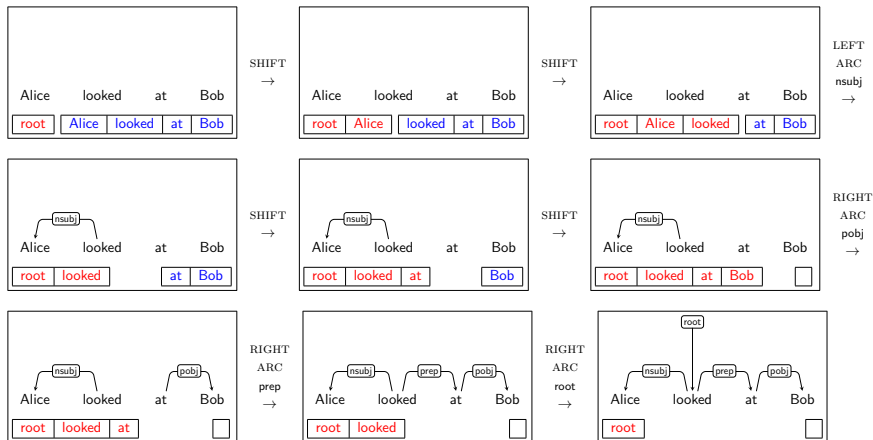
Projective tree (no crossing arcs \Leftrightarrow all sub-trees are sub-strings):




Non-projective tree (cannot be parsed by arc-standard):



Another example for arc-standard transition sequence



Another example for arc-standard transition sequence



Might be a good idea to attach looked → at as soon as possible?

Arc-eager transition system (Nivre, 2004)

SHIFT (same)	move one item from the buffer to the stack: $(\Sigma, i B, A) \Rightarrow (\Sigma i, B, A)$
LEFT-ARC $_{\ell}$	create arc $b_0 \rightarrow s_0$ with label $\ell \in \mathcal{L}$ and remove s_0 : $(\Sigma i, j B, A) \Rightarrow (\Sigma, j B, A \cup \{(j, \ell, i)\})$ Condition: $i \neq 0$ and i has no head
RIGHT-ARC $_{\ell}$	create arc $s_0 \rightarrow b_0$ with label $\ell \in \mathcal{L}$ and shift b_0 : $(\Sigma i, j B, A) \Rightarrow (\Sigma i j, B, A \cup \{(i, \ell, j)\})$
REDUCE	remove s_0 : $(\Sigma i, B, A) \Rightarrow (\Sigma, B, A)$ Condition: i has a head

Arc-eager transition system (Nivre, 2004)

SHIFT (same)	move one item from the buffer to the stack: $(\Sigma, i B, A) \Rightarrow (\Sigma i, B, A)$
LEFT-ARC $_{\ell}$	create arc $b_0 \rightarrow s_0$ with label $\ell \in \mathcal{L}$ and remove s_0 : $(\Sigma i, j B, A) \Rightarrow (\Sigma, j B, A \cup \{(j, \ell, i)\})$ Condition: $i \neq 0$ and i has no head
RIGHT-ARC $_{\ell}$	create arc $s_0 \rightarrow b_0$ with label $\ell \in \mathcal{L}$ and shift b_0 : $(\Sigma i, j B, A) \Rightarrow (\Sigma i j, B, A \cup \{(i, \ell, j)\})$
REDUCE	remove s_0 : $(\Sigma i, B, A) \Rightarrow (\Sigma, B, A)$ Condition: i has a head

Initial configuration: same as arc-standard.

Terminal configuration (Σ does not have to be $[0]$):

$$c_t = (\Sigma, [], A)$$

Properties of the arc-eager system

Soundness and completeness are the same as arc-standard.


Complexity: **at most** $2n$.

Properties of the arc-eager system

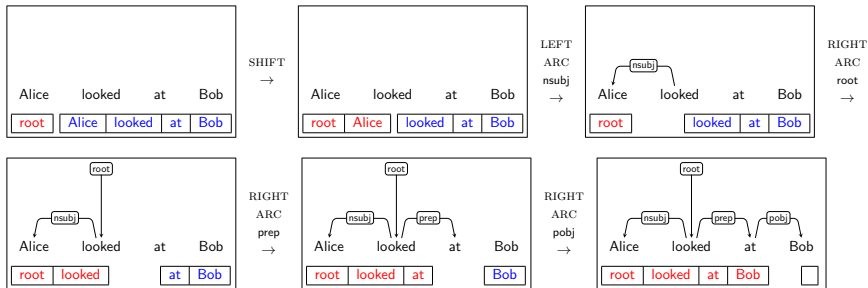
Soundness and completeness are the same as arc-standard.

Complexity: **at most** $2n$.

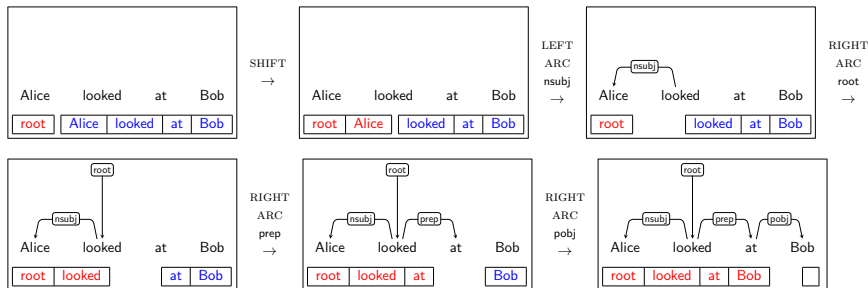
Builds left-dependents bottom-up and right-dependents top-down.
Increased incrementality—no need to wait for the whole sub-tree to be complete before attaching it.



Example arc-eager transition sequence



Example arc-eager transition sequence



Shorter than $2n$ since we can skip the final REDUCE transitions.

Greedy transition-based parsing

Transition-based (shift-reduce) parsing

To actually parse text, we need to decide which transitions to take.

$$P(t_1, \dots, t_m | w) = \prod_{i=1}^m P(t_i | t_1, \dots, t_{i-1}, w) = \prod_{i=1}^m P(t_i | c_{i-1})$$

so inference is

$$\arg \max_{t_1, \dots, t_m \in \mathcal{T}} \prod_{i=1}^m P(t_i | c_{i-1})$$

But training examples are trees, not sequences.

To learn this score, we need an *oracle* to tell the correct sequence:

$$o(T) = (t_1, \dots, t_m)$$

Transition-based (shift-reduce) parsing

To actually parse text, we need to decide which transitions to take.

$$P(t_1, \dots, t_m | w) = \prod_{i=1}^m P(t_i | t_1, \dots, t_{i-1}, w) = \prod_{i=1}^m P(t_i | c_{i-1})$$


so inference is

$$\arg \max_{t_1, \dots, t_m \in \mathcal{T}} \prod_{i=1}^m P(t_i | c_{i-1})$$

But training examples are trees, not sequences.

To learn this score, we need an *oracle* to tell the correct sequence:

$$o(T) = (t_1, \dots, t_m)$$



⇒ SHIFT, SHIFT, LEFT-ARC_{nsubj}, SHIFT, RIGHT-ARC_{dobj}, RIGHT-ARC_{root}

Oracle for arc-standard

```
while  $B \neq []$  and  $\Sigma \neq [0]$  do  
  if  $s_0 \xrightarrow{\ell} s_1$  and  $s_1$  has all its children and  $s_1 \neq 0$  then  
    return LEFT-ARC $_{\ell}$   
  else if  $s_1 \xrightarrow{\ell} s_0$  and  $s_0$  has all its children and  $s_0 \neq 0$  then  
    return RIGHT-ARC $_{\ell}$   
  else  
    return SHIFT  
  end if  
end while
```

Oracle for arc-eager

```
while  $B \neq []$  do  
  if  $b_0 \xrightarrow{\ell} s_0$  then  
    return LEFT-ARC $_{\ell}$   
  else if  $s_0 \xrightarrow{\ell} b_0$  then  
    return RIGHT-ARC $_{\ell}$   
  else if  $s_0$  has all its children and a head then  
    return REDUCE  
  else  
    return SHIFT  
  end if  
end while
```

Greedy transition-based parsing

In *greedy parsing*, instead of

$$(t_1, \dots, t_m) = \arg \max_{t'_1, \dots, t'_m \in \mathcal{T}} \prod_{i=1}^m P(t'_i | c_{i-1})$$

we select each transition separately and sequentially:

$$t_i = \arg \max_{t'_i \in \mathcal{T}} P(t'_i | c_{i-1}) \quad i = 1, \dots, m$$

Greedy transition-based parsing

In *greedy parsing*, instead of

$$(t_1, \dots, t_m) = \arg \max_{t'_1, \dots, t'_m \in \mathcal{T}} \prod_{i=1}^m P(t'_i | c_{i-1})$$

we select each transition separately and sequentially:

$$t_i = \arg \max_{t'_i \in \mathcal{T}} P(t'_i | c_{i-1}) \quad i = 1, \dots, m$$

A score $s(t, c)$ estimates this probability. Parsing algorithm:

$c \leftarrow c_s(w)$

while $c \notin C_t$ **do**

$c \leftarrow \left(\arg \max_{t \in \mathcal{T}} s(t, c) \right)(c)$

end while

Transition classifiers

Learn the score giving maximum probability to oracle transitions:

$$\arg \max_{s \in \mathcal{S}} \sum_{i=1}^m s(t_i^*, c_{i-1}^*)$$

where t_1^*, \dots, t_m^* (and c_1^*, \dots, c_m^*) are determined by the oracle.

Transition classifiers

Learn the score giving maximum probability to oracle transitions:

$$\arg \max_{s \in \mathcal{S}} \sum_{i=1}^m s(t_i^*, c_{i-1}^*)$$

where t_1^*, \dots, t_m^* (and c_1^*, \dots, c_m^*) are determined by the oracle.

Possible hypothesis classes \mathcal{S} :

- 1 Linear (perceptron)
- 2 Feedforward neural networks
- 3 Recurrent neural networks (RNN, LSTM, GRU)

And others, e.g. SVM.

Linear transition classifier (Nivre, 2003)

Given features $\mathbf{f} = (f_1, \dots, f_K) : \mathcal{C} \rightarrow \mathbb{R}^K$, learn weights $W_{|\mathcal{T}| \times K}$:

$$s(t, c) = [W \cdot \mathbf{f}(c)]_t$$

Typically trained by the perceptron algorithm, with binary features: words, POS and existing arc labels of stack and buffer nodes, and their heads and dependents.

NN transition classifier (Chen and Manning, 2014)

Dense **embedding** features instead of sparse binary features.
Trained with backpropagation and stochastic gradient descent.

Feedforward NN architecture:


Softmax layer:

$$p = \text{softmax}(W_2 h)$$

Hidden layer:


$$h = (W_1^w x^w + W_1^t x^t + W_1^l x^l + b_1)^3$$

Input layer: $[x^w, x^t, x^l]$




Bi-RNN encoder (Kiperwasser and Goldberg, 2016)

Deep bidirectional LSTM RNN for input representations.



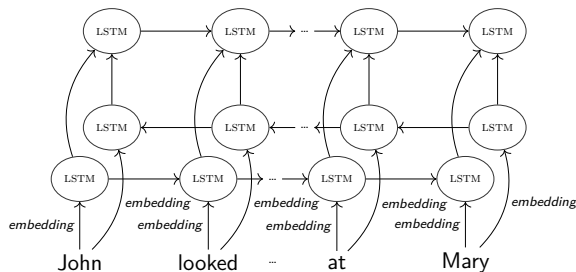
Bi-RNN encoder (Kiperwasser and Goldberg, 2016)

Deep bidirectional LSTM RNN for input representations.




Bi-RNN encoder (Kiperwasser and Goldberg, 2016)

Deep bidirectional LSTM RNN for input representations.



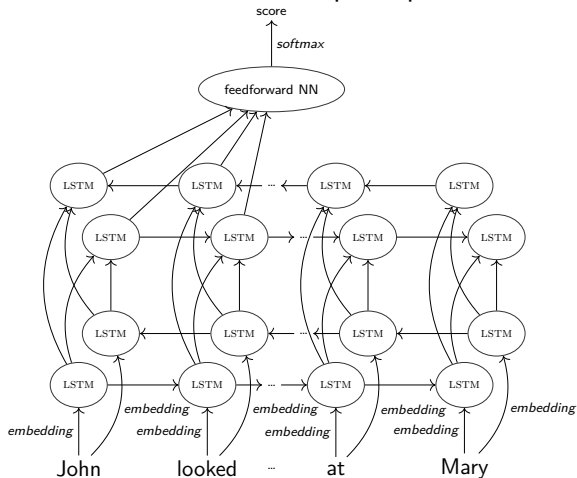
Bi-RNN encoder (Kiperwasser and Goldberg, 2016)

Deep bidirectional LSTM RNN for input representations.



Bi-RNN encoder (Kiperwasser and Goldberg, 2016)

Deep bidirectional LSTM RNN for input representations.



Empirical comparison

Evaluation on PTB-SD¹:

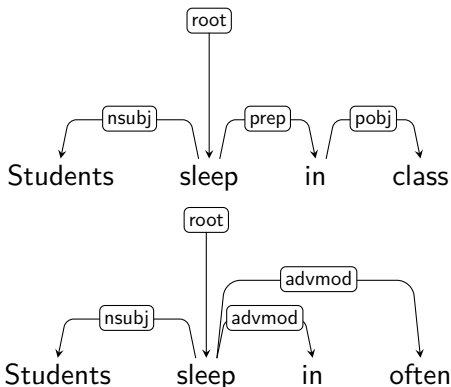
	UAS	LAS
Graph-based		
MSTParser	90.7	87.6
Transition-based		
Linear (Zhang and Nivre, 2011)	89.6	87.4
Feedforward NN (Chen and Manning, 2014)	91.8	89.6
BiLSTM (Kiperwasser and Goldberg, 2016)	93.9	91.9

¹Penn Treebank Wall-Street Journal (WSJ) with Stanford Dependencies

Dealing with error propagation

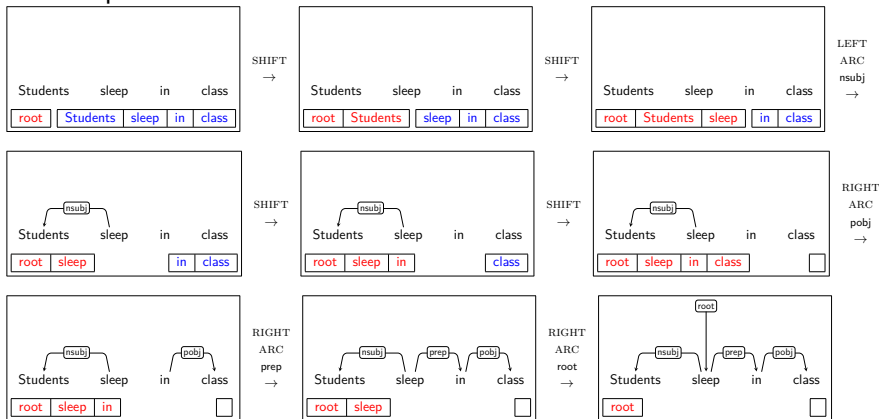
Error propagation

Greedy transition-based parsers do not recover well from errors.




Error propagation example

Correct parse:




Error propagation example

Error during parse:



Error propagation example

Error during parse:



Results in a state never seen during training.

Solutions for error propagation

- Better transition classifier with context "look-ahead" (LSTM).
- Beam search and structured training.
- Dynamic oracle and training with exploration.

Solutions for error propagation

- Better transition classifier with context "look-ahead" (LSTM).
- **Beam search and structured training.**
- Dynamic oracle and training with exploration.

Beam search and structured training

Reminder—greedy parsing algorithm:

```
 $c \leftarrow c_s(w)$   
while  $c \notin C_t$  do  
   $c \leftarrow \left( \arg \max_{t \in \mathcal{T}} s(t, c) \right)(c)$   
end while
```

Beam search and structured training

Reminder—greedy parsing algorithm:

```
 $c \leftarrow c_s(w)$   
while  $c \notin C_t$  do  
   $c \leftarrow \left( \arg \max_{t \in \mathcal{T}} s(t, c) \right)(c)$   
end while
```

With *beam search*, we instead keep the k best transition sequences where k is the beam size.

Beam search and structured training

Reminder—greedy parsing algorithm:

```
 $c \leftarrow c_s(w)$   
while  $c \notin C_t$  do  
   $c \leftarrow \left( \arg \max_{t \in \mathcal{T}} s(t, c) \right)(c)$   
end while
```

With *beam search*, we instead keep the k best transition sequences where k is the beam size.

$k = 1$ is greedy parsing.

Beam search algorithm

Maintain beam Q of top-scoring configurations with their scores:

$$Q \leftarrow \left\{ (c_s(w), s) \right\}$$

while there exists $(c, s) \in Q$ s.t. $c \notin C_t$ **do**

$$Q \leftarrow \text{SELECT} \left(k, \left\{ (t(c), s + s(t, c)) \mid (c, s) \in Q, t \in \mathcal{T} \right\} \right)$$

end while

return $\text{SELECT}(1, Q)$

Beam search algorithm

Maintain beam Q of top-scoring configurations with their scores:

$$Q \leftarrow \left\{ \left(c_s(w), s \right) \right\}$$

while there exists $(c, s) \in Q$ s.t. $c \notin C_t$ **do**

$$Q \leftarrow \text{SELECT} \left(k, \left\{ \left(t(c), s + s(t, c) \right) \mid (c, s) \in Q, t \in \mathcal{T} \right\} \right)$$

end while

return $\text{SELECT}(1, Q)$

If the top sequence has an error, a lower-scoring one might be better in the long run.

Empirical comparison

Evaluation on PTB-SD:

	k	UAS	LAS
Greedy			
Linear (Zhang and Nivre, 2011)	1	89.6	87.4
Feedforward NN (Chen and Manning, 2014)	1	91.8	89.6
Beam search			
Linear (Bohnet and Nivre, 2012)	40	93.2	91.1
Feedforward NN+perceptron (Weiss et al., 2015)	8	93.9	92
Dynamic oracle			
BiLSTM (Kiperwasser and Goldberg, 2016)	1	93.9	91.9

References I

Bohnet, B. and Nivre, J. (2012).

A transition-based system for joint part-of-speech tagging and labeled non-projective dependency parsing.
In *Proc. of EMNLP-CoNLL*, pages 1455–1465.

Chen, D. and Manning, C. (2014).

A fast and accurate dependency parser using neural networks.
In *Proc. of EMNLP*, pages 740–750.

Kiperwasser, E. and Goldberg, Y. (2016).

Simple and accurate dependency parsing using bidirectional LSTM feature representations.
TACL, 4:313–327.

Nivre, J. (2003).

An efficient algorithm for projective dependency parsing.
In *Proc. of IWPT*, pages 149–160.

Nivre, J. (2004).

Incrementality in deterministic dependency parsing.
In Keller, F., Clark, S., Crocker, M., and Steedman, M., editors, *Proceedings of the ACL Workshop Incremental Parsing: Bringing Engineering and Cognition Together*, pages 50–57, Barcelona, Spain. Association for Computational Linguistics.

Weiss, D., Alberti, C., Collins, M., and Petrov, S. (2015).

Structured training for neural network transition-based parsing.
arXiv preprint arXiv:1506.06158.

References II

Zhang, Y. and Nivre, J. (2011).

Transition-based dependency parsing with rich non-local features.

In *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies*, pages 188–193.